

Attentive TCN-WGAN-GP for Low-False-Positive and Explainable Network Intrusion Detection

Jie Li

School of Software, Beihang University, Beijing, China

Email: zb2557106@buaa.edu.cn

Abstract—Network intrusion detection is useful in practice only when alarms are timely, low-noise, and auditable by security operators. Existing flow-level classifiers often ignore temporal context across adjacent flows, while many deep anomaly detectors report aggregate metrics without clarifying whether their detection granularity, false-positive behavior, and latency are suitable for operational monitoring. This paper presents Attentive TCN-WGAN-GP, an attention-enhanced Temporal Convolutional Network based Wasserstein Generative Adversarial Network with Gradient Penalty, for low-false-positive window-level intrusion detection. The proposed approach converts network flows into sliding windows, models benign temporal patterns using a TCN-based generator and critic, stabilizes adversarial training with a WGAN-GP objective, and computes a fused anomaly score from critic evidence and feature-deviation evidence. We evaluate the detector on CICIDS2017, SWaT, UNSW-NB15, and TON_IoT against classical, deep, and modern anomaly-detection baselines including GANomaly, DeepSVDD, and TranAD. On CICIDS2017, Attentive TCN-WGAN-GP achieves the best low-FPR calibrated F1 among the compared methods, with 0.7958 F1 and 0.0036 observed benign test FPR at target FPR 0.05, and reaches 0.8783 F1 and 0.9678 recall at target FPR 0.15. On SWaT, it obtains 0.9950 F1 and zero observed benign test FPR under a formal anomaly-detection split, supporting transfer to an industrial-control setting. On UNSW-NB15, it remains competitive under a higher-recall operating point but does not dominate all baselines, highlighting the difficulty of threshold transfer under distribution shift. TON_IoT is treated as an out-of-domain limitation case rather than a primary win condition. We further provide a multi-view XAI case study that combines attention weights, Integrated Gradients, and SHAP attributions to audit a single alarm window.

Index Terms—network intrusion detection, low false positives, temporal convolutional network, generative adversarial network, explainable AI

I. INTRODUCTION

Network intrusion detection systems must identify malicious behavior from high-volume, evolving traffic while keeping false alarms low enough for deployment. A useful detector for security operations should reason over short temporal contexts, support online scoring, and provide evidence that helps analysts understand why a window of traffic was flagged.

Many machine-learning intrusion detection pipelines operate at the flow level: each flow is treated as an independent sample, and the model predicts whether that flow is benign or malicious. This formulation is simple and useful, but it can miss temporal patterns across neighboring flows and can blur runtime claims, because a strong flow-level classifier is not automatically a stable streaming detector. For a commu-

nications and network security setting, the unit of analysis should therefore match the unit of response: an alarm should summarize a recent traffic segment rather than a single isolated record.

The false-alarm constraint is especially important in operational security monitoring. A detector with high recall but excessive benign false positives may overwhelm analysts and reduce trust in alarms. For this reason, we evaluate not only ranking metrics such as AUC and average precision, but also calibrated operating points under explicit false-positive-rate budgets.

This paper studies a window-level alternative. We convert consecutive network flows into fixed-length windows and train a TCN-GAN model to capture benign temporal structure. To improve stability and evidence aggregation, we upgrade the basic TCN-GAN with WGAN-GP, short for Wasserstein Generative Adversarial Network with Gradient Penalty, and an attention-pooling discriminator. We call the resulting framework Attentive TCN-WGAN-GP. At inference time, the anomaly score fuses the discriminator’s unknownness signal with feature-space deviation from benign reference windows.

Our target claim is deliberately narrow and testable: after fixing the window size, stride, and score-fusion ratio, Attentive TCN-WGAN-GP should outperform same-granularity baselines at low-noise operating points and provide meaningful post-hoc explanations for alarmed windows. We do not claim global hyperparameter optimality or universal dominance across all traffic domains.

This paper makes the following contributions:

- We propose Attentive TCN-WGAN-GP, an attentive TCN-WGAN-GP pipeline for low-false-positive, window-level network intrusion detection.
- We formulate the detector around an operational monitoring workflow: sliding-window scoring, benign-only threshold calibration, explicit false-positive-rate budgets, and throughput reporting.
- We combine attention pooling, WGAN-GP training, and fused critic/feature-deviation scoring in a single detection framework, and isolate these choices through ablation.
- We validate the detector beyond CICIDS2017 using SWaT, UNSW-NB15, and TON_IoT, separating positive transfer, stress-test behavior under distribution shift, and explicit out-of-domain limitations.
- We add alarm-audit XAI analysis based on attention weights, Integrated Gradients, and SHAP to explain tem-

poral and feature-level evidence behind single-window anomaly decisions.

II. BACKGROUND AND MOTIVATION

A. Operational Monitoring Goal

We consider a network monitoring setting in which a defender observes flow-level telemetry from an enterprise, cloud, or industrial network and must raise timely alarms for suspicious traffic segments. The detector is trained primarily on benign traffic, which reflects the common deployment assumption that new attack variants may not be available during training. Attack labels are used for evaluation, model selection analysis, and comparison with baselines, but the proposed adversarial detector learns the benign-window distribution.

The security goal is not to assign a perfect semantic attack label to every flow. Instead, the goal is to produce a low-noise alarm stream that prioritizes windows likely to contain malicious behavior. This framing makes three requirements central: the alarm unit should preserve short-term traffic context, the threshold should be calibrated under an explicit false-positive budget, and the final alarm should be inspectable enough for an analyst to decide whether escalation is warranted.

B. Window-Level Detection

Let a network trace be represented as an ordered sequence of flow feature vectors. A flow-level detector maps each vector to a label independently. In contrast, a window-level detector maps a sequence window to a score:

$$X_i = [x_i, x_{i+1}, \dots, x_{i+W-1}], \quad s_i = f(X_i), \quad (1)$$

where W is the window size and consecutive windows are separated by stride r . This formulation is better aligned with real-time monitoring because the detector can emit an alarm whenever the score of the current window exceeds a calibrated threshold.

The choice of W and r controls a trade-off. Larger windows contain more temporal context but increase latency and memory use; smaller strides improve temporal resolution but create more windows and higher compute cost. Our experiments therefore report both detection quality and throughput.

C. Temporal Convolutional Networks and GANs

Temporal convolutional networks use causal or dilated one-dimensional convolutions to capture sequence patterns efficiently [1]. GANs train a generator and discriminator in an adversarial objective [2]; for anomaly detection, the discriminator can learn a boundary around normal data. Because vanilla GAN training may be unstable, Wasserstein GAN (WGAN) objectives and WGAN-GP, namely Wasserstein Generative Adversarial Network with Gradient Penalty, are often used to improve critic training [3], [4]. The Wasserstein objective replaces binary real/fake classification with a critic score that approximates the distance between the real and generated distributions, while the gradient penalty regularizes the critic to satisfy the Lipschitz constraint required by the Wasserstein formulation.

III. RELATED WORK

A. Network Intrusion and Anomaly Detection

Classical intrusion detection pipelines often use flow-level features and supervised or one-class models. Earlier benchmarks such as KDD Cup 1999 and NSL-KDD helped standardize IDS evaluation, but their traffic distributions and attack patterns are widely considered limited for modern network settings [5]. Later datasets such as UNSW-NB15 and CICIDS2017 introduced more diverse flow features and attack scenarios [6], [7]. CICIDS2017 remains a common benchmark because it contains benign traffic and several attack families collected across multiple days.

Traditional machine-learning detectors include tree ensembles, support-vector methods, nearest-neighbor methods, and one-class anomaly detectors. Isolation Forest isolates anomalous points using random partitioning [8], while One-Class SVM estimates a boundary around normal samples [9]. These methods are useful baselines because they are simple and strong on tabular flow features. However, independent flow classification can miss temporal evidence that emerges across adjacent flows, and threshold calibration can strongly affect recall and false positives.

Unsupervised and online anomaly detection methods reduce dependence on attack labels. Kitsune, for example, uses an ensemble of autoencoders for online network intrusion detection and emphasizes deployability on resource-constrained devices [10]. Recent work has also explored Transformer-based NIDS frameworks and cloud-security intrusion detectors, showing that attention architectures are increasingly relevant for flow-based security analytics [11]–[13]. A particularly related direction combines GANs, TCNs, and self-attention for unsupervised intrusion detection [14]; recent IP-flow GAN work also evaluates TCN, LSTM, and CNN variants with SHAP-based analysis on CIC-style traffic datasets [15]. Our work shares the goal of online anomaly detection, but differs in its explicit low-FPR calibration protocol, WGAN-GP ablation, fused anomaly scoring, and multi-view XAI case study.

B. Deep Sequence Models for Anomaly Detection

Deep anomaly detection has used recurrent, convolutional, forecasting-based, Transformer-based, and autoencoder-based models to capture temporal or reconstruction patterns. Recurrent models have been used for time-series anomaly detection and log anomaly detection, where the model learns expected sequential behavior and flags deviations [16], [17]. Autoencoder variants reconstruct normal inputs and treat high reconstruction error as evidence of abnormality [18], [19]. DeepSVDD trains a deep representation using a one-class objective [20], while TranAD uses Transformer encoders for multivariate time-series anomaly detection [21]. Recent general time-series models such as Anomaly Transformer and TimesNet further show the importance of attention, association, and temporal variation modeling for unsupervised anomaly detection [22], [23]. A recent review of time-series analysis for cyber security further highlights intrusion detection as a key application area and emphasizes challenges

such as heterogeneous feature types, sparsity, and temporal aggregation [24]. Industrial IoT anomaly-detection work has also moved toward dynamic and adaptive sequence models for multivariate telemetry [25]. These approaches are natural for sequence data, but recurrent or Transformer evaluation can be slower, and reconstruction quality does not always align with detection quality.

Temporal convolutional networks offer a practical alternative because dilated convolutions can capture long-range sequence patterns with efficient parallel computation [1]. Compared with recurrent models, TCNs are easier to batch across windows and can provide predictable inference latency. Our method uses TCN blocks inside an adversarial detector and reports window-level throughput to keep the real-time monitoring claim grounded.

C. GAN-Based Detection and Training Stability

GAN-based anomaly detection uses adversarial training to learn a normal-data manifold, a discriminator boundary, or an embedding where abnormal samples are easier to separate. AnoGAN showed that GANs can support anomaly detection by comparing samples with generated normal patterns [26], while later variants such as f-AnoGAN improved inference efficiency [27]. GANomaly introduced an encoder-decoder-encoder adversarial detector that has become a common deep anomaly baseline [28]. However, vanilla GAN objectives can be unstable, especially under imbalanced or limited-data settings. WGAN and WGAN-GP improve training by replacing the discriminator classification objective with a Wasserstein critic and gradient penalty [3], [4]. Our ablation explicitly evaluates whether WGAN-GP improves a TCN-based adversarial detector when the window size, stride, and score fusion are held fixed.

D. Explainability and Attention

XAI is increasingly important for security systems because analysts need to understand why an alarm was raised, not only whether it was raised. Feature-attribution methods such as saliency maps, Integrated Gradients, LIME, and SHAP provide post-hoc explanations of model predictions [29]–[32]. Recent IDS studies have increasingly combined attention models with SHAP, LIME, drift awareness, and cross-dataset evaluation to improve analyst trust [33], [34]. Attention weights can also expose where a model aggregates evidence, but attention should not automatically be treated as a faithful explanation of the final prediction [35]. For this reason, we report both attention weights and attribution methods: attention describes temporal aggregation inside the discriminator, while Integrated Gradients and SHAP estimate which time steps and traffic features contribute to the final anomaly score.

IV. METHOD

A. Data Representation

Given numeric flow features, we clean and scale each feature, preserve the temporal order within each file split, and construct sliding windows of length W . A window is labeled

Overall Framework of Attentive TCN-WGAN-GP




Fig. 1. Overall framework of the proposed Attentive TCN-WGAN-GP pipeline.

Fig. 1. End-to-end workflow of the proposed window-level intrusion detection framework. Flow records are cleaned, scaled, converted into sliding windows, scored by the trained detector, calibrated under a benign false-positive budget, and audited with post-hoc XAI evidence.

anomalous if any flow inside the window is labeled as an attack. This produces the window-level evaluation target used throughout the main experiments.

This representation is intentionally different from a flow-level classifier. The input to the detector is not a single flow vector, but a short ordered traffic segment. The label of a window is used only for evaluation and threshold analysis; the adversarial model itself is trained to capture the benign-window distribution. As traffic arrives, the system advances by stride r , forms the newest window, computes an anomaly score, and compares it with a calibrated threshold.

B. Modeling the Detector Upgrade

We name the proposed framework Attentive TCN-WGAN-GP to make its components explicit. “Attentive” refers to the discriminator’s temporal attention pooling, “TCN” refers to the Temporal Convolutional Network backbone used for window modeling, and “WGAN-GP” refers to Wasserstein Generative Adversarial Network with Gradient Penalty, the adversarial objective used for stable training. The name therefore describes both the architectural modification and the training objective.

The proposed model can be viewed as an upgrade from a basic TCN-GAN detector to an attentive TCN-WGAN-GP detector with fused anomaly scoring. Let $X \in \mathbb{R}^{W \times d}$ denote a traffic window with W time steps and d flow features. The generator G_θ maps latent noise $z \sim p(z)$ to a synthetic benign-like window,

$$\tilde{X} = G_\theta(z), \quad \tilde{X} \in \mathbb{R}^{W \times d}, \quad (2)$$

and the discriminator or critic D_ϕ maps a real or generated window to a scalar evidence score. The basic TCN-GAN version uses TCN encoders in both G_θ and D_ϕ and trains them with a standard adversarial real/fake objective:

$$\begin{aligned} \mathcal{L}_D^{\text{vanilla}} = & -\mathbb{E}_{X \sim P_r} [\log \sigma(D_\phi(X))] \\ & -\mathbb{E}_{z \sim p(z)} [\log(1 - \sigma(D_\phi(G_\theta(z))))], \end{aligned} \quad (3)$$

$$\mathcal{L}_G^{\text{vanilla}} = -\mathbb{E}_{z \sim p(z)} [\log \sigma(D_\phi(G_\theta(z)))]. \quad (4)$$




Fig. 2. Architecture of the proposed Attentive TCN-WGAN-GP model.

Fig. 2. Architecture of Attentive TCN-WGAN-GP. A TCN generator synthesizes benign-like windows, while a TCN critic with attention pooling scores real and generated windows. WGAN-GP stabilizes adversarial training, and fused scoring combines critic and representation-deviation evidence at inference time.

This baseline gives a temporal adversarial detector, but it has two weaknesses: the discriminator aggregates all time steps uniformly, and the vanilla GAN objective can be unstable. We therefore model the final detector as three explicit modifications:

$$\begin{aligned} M_0 &= \text{TCN-GAN}, \\ M_1 &= M_0 + \text{attention pooling}, \\ M_2 &= M_1 + \text{WGAN-GP}, \\ S(X) &= \text{fused scoring}(M_2, X). \end{aligned} \quad (5)$$

The first modification changes temporal aggregation, the second changes the adversarial training objective, and the third changes the inference-time anomaly score. This formulation makes the model upgrade measurable in ablation.

C. Attentive TCN-WGAN-GP Architecture

The generator receives a latent vector z and maps it through a linear layer followed by TCN blocks and a 1×1 convolution to produce a synthetic window. The discriminator maps an input window through TCN blocks to obtain temporal hidden features. A time-wise classifier produces logits for each time step. With mean pooling, the discriminator averages these logits. With attention pooling, the discriminator learns a normalized weight over time steps and aggregates the logits by weighted sum:

$$D(X) = \sum_{t=1}^W a_t \ell_t, \quad a_t = \frac{\exp(q_t)}{\sum_{j=1}^W \exp(q_j)}. \quad (6)$$

Here ℓ_t is the discriminator logit at time t , and a_t is the attention weight.

The attention module is not introduced as a stand-alone explanation method. It is a learnable pooling mechanism inside the discriminator. Mean pooling assumes that all positions in the window contribute equally to the discriminator score, while attention pooling allows the model to emphasize positions that provide stronger evidence about whether the whole window resembles benign traffic.

D. WGAN-GP Training

WGAN-GP stands for Wasserstein Generative Adversarial Network with Gradient Penalty. In this variant, the discriminator is treated as a critic rather than a probability classifier. The critic assigns larger scores to real benign windows and smaller scores to generated windows, and the gradient penalty constrains the critic to change smoothly between real and generated samples. The critic minimizes:

$$\begin{aligned} \mathcal{L}_D &= \mathbb{E}_{\tilde{X} \sim P_g} [D(\tilde{X})] - \mathbb{E}_{X \sim P_r} [D(X)] \\ &\quad + \lambda \mathbb{E}_{\tilde{X}} [\|\nabla_{\tilde{X}} D(\tilde{X})\|_2 - 1]^2, \end{aligned} \quad (7)$$

while the generator minimizes:

$$\mathcal{L}_G = -\mathbb{E}_{\tilde{X} \sim P_g} [D(\tilde{X})]. \quad (8)$$

The gradient penalty coefficient is $\lambda = 10$ in the current main script, and the critic is updated five times per generator update.

E. Fused Anomaly Scoring

The discriminator score captures how real the window appears to the discriminator. We also compute a feature-deviation score by embedding windows with discriminator features and measuring the distance to a benign reference distribution. Both scores are min-max normalized with benign reference statistics and fused:

$$S(X) = \alpha S_D(X) + (1 - \alpha) S_F(X), \quad (9)$$

where S_D is the discriminator unknownness score and S_F is the feature-deviation score. The final threshold is calibrated on benign windows to target a fixed false positive rate.

This fused score is useful because the critic and the feature embedding capture related but not identical evidence. The critic score reflects the adversarial realism signal learned during GAN training, whereas the feature-deviation score compares the hidden representation of a test window against benign reference statistics. The parameter α controls this trade-off.

F. Explainability

For post-hoc explanation, we compute gradient-times-input and Integrated Gradients style attributions over the final anomaly score:

$$A_{t,k} = \left| X_{t,k} \frac{\partial S(X)}{\partial X_{t,k}} \right|. \quad (10)$$

We summarize attributions over time and over features separately, and compare benign and anomalous windows. When attention pooling is enabled, we also report the learned temporal attention weights. For a single-window case study, we additionally compute SHAP GradientExplainer attributions with a small benign background set. These explanations are not used to train the current model; they are used to audit whether alarms are driven by plausible temporal regions and features. Attention is therefore treated as an internal evidence aggregation signal, not as a faithful explanation by itself.

V. EXPERIMENTAL SETUP

A. Datasets and Splits

We use CICIDS2017 as the main benchmark [7]. The training split uses Tuesday-WorkingHours, Wednesday-workingHours, Thursday-WorkingHours-Morning-WebAttacks, and Thursday-WorkingHours-Afternoon-Infiltration. The test split uses Friday-WorkingHours-Morning, Friday-WorkingHours-Afternoon-DDos, and Friday-WorkingHours-Afternoon-PortScan. The split is chronological by file rather than random over all windows, which reduces leakage of day-specific traffic conditions.

To test whether the detector generalizes beyond the main benchmark, we add three cross-dataset evaluations. SWaT is used as an industrial-control validation setting [36], where benign data are split into train, calibration, and benign-test segments, and the attack file is used as the anomalous test segment. Recent SWaT-focused work continues to use autoencoder-style detectors for ICS anomaly detection, reinforcing SWaT as an active industrial-control benchmark [37]. UNSW-NB15 [6] is used as a third authoritative IDS dataset with its official training and testing CSV files; we use numeric features and remove identifiers and categorical labels. TON_IoT [38] is treated as a cross-domain telemetry challenge set rather than as a primary win condition, because its Linux telemetry stream differs substantially from flow and process-control traffic. Recent work also evaluates hybrid IDS models on UNSW-NB15 and TON_IoT, reinforcing their role as useful but challenging cross-dataset benchmarks [39].

All methods are evaluated at the window level with stride $r = 16$ and anomaly ratio threshold 0.15. The main proposed model and in-repository baselines use $W = 128$. External sequence baselines that impose or conventionally use a different input length are reported with their window length explicitly shown in the result tables. A window is anomalous if at least 15% of its contained records are labeled as attacks. This keeps the evaluation granularity visible and avoids mixing flow-level and window-level claims.

B. Baselines

The main comparison uses same-granularity window-level baselines: IsolationForest, OneClassSVM, Random Forest, MLP, LSTM-AE, and LSTM-AD. We further add modern anomaly-detection baselines, including GANomaly [28], DeepSVDD [20], a compact TranAD-style Transformer reconstruction model [21], Anomaly Transformer [22], TimesNet [23], and ALoRa [40]. We additionally ran Time-Series-Library family variants, DLinear [41] and Autoformer [42], under the same window-level protocol as supplementary external references, following the Time-Series-Library benchmark line [23], [43]. Legacy flow-level results are not mixed into the main comparison because their sample granularity differs from the proposed window-level detector.

The baseline set covers four groups. IsolationForest and OneClassSVM represent classical unsupervised anomaly detection. Random Forest and MLP represent supervised

window-level classifiers trained with window labels. LSTM-AE and LSTM-AD represent deep sequence baselines based on reconstruction or sequence modeling. GANomaly, DeepSVDD, TranAD, Anomaly Transformer, TimesNet, ALoRa, and supplementary Time-Series-Library variants (DLinear and Autoformer) represent modern deep anomaly-detection baselines. This grouping helps separate the benefit of temporal benign-pattern modeling from the effects of supervision, reconstruction, and generic deep anomaly scoring.

C. Metrics

We report AUC, average precision (AP), calibrated F1, calibrated precision, calibrated recall, calibrated false positive rate, and inference throughput. Calibrated metrics use thresholds chosen on benign training or calibration windows. We report a strict low-FPR operating point at target FPR 0.05 and also analyze a higher-recall operating point at target FPR 0.15. Because train and test benign distributions may shift, we report the observed test benign FPR separately from the requested target FPR.

This distinction is important for network security evaluation. A threshold that satisfies a requested FPR on benign calibration data may still produce more benign alarms after traffic distribution shift. Reporting the observed benign test FPR makes this failure mode visible instead of hiding it behind a single F1 score. Throughout the results, we therefore interpret high recall together with the benign-alarm cost required to obtain it.

D. Implementation Environment

Experiments were run on an Apple M3 MacBook Pro with 8 CPU cores and 8 GB unified memory, using Python 3.10.20 and PyTorch 2.2.2 on CPU. Reported throughput is therefore CPU-only and should be interpreted as an offline feasibility measurement rather than a tuned gateway benchmark.

E. Hyperparameter Selection

We use staged hyperparameter selection. First, existing window-sweep results motivate $W = 128$ as a context/latency trade-off. Second, with W fixed, the final script sweeps stride values and selects the best stride by calibrated F1. Third, with window and stride fixed, it sweeps the score-fusion parameter α . The selected configuration is written to `selected_main_config.json` and reused for baseline comparison and ablation. This procedure is not a global search; it is designed to prevent moving baselines during ablation.

This staged protocol is a deliberate fairness constraint. Without freezing W , r , and α , the ablation could accidentally compare different data granularities or different thresholding behavior rather than the actual architectural changes. After the configuration is frozen, the baseline comparison and the ablation use the same window size, stride, anomaly-ratio threshold, score mode, and target-FPR protocol.

TABLE I
WINDOW-LEVEL COMPARISON AT THE STRICT LOW-FPR OPERATING POINT (TARGET FPR 0.05).

Model	AUC	AP	F1	Recall	Test FPR
Attentive TCN-WGAN-GP	0.9464	0.9351	0.7958	0.6640	0.0036
MLP	0.7537	0.7386	0.7438	0.6900	0.1210
LSTM-AE	0.4262	0.4201	0.5943	1.0000	1.0000
LSTM-AD	0.5867	0.4705	0.5943	1.0000	1.0000
RF	0.6623	0.7126	0.5578	0.4763	0.1695
OneClassSVM	0.8139	0.6506	0.2605	0.1675	0.0870
IsolationForest	0.3961	0.4525	0.2482	0.1533	0.0601

TABLE II
CICIDS2017 COMPARISON WITH MODERN DEEP ANOMALY-DETECTION BASELINES AT TARGET FPR 0.05.

Model	AUC	AP	F1	Recall	Test FPR
Attentive TCN-WGAN-GP	0.9464	0.9351	0.7958	0.6640	0.0036
GANomaly	0.7409	0.7706	0.7223	0.7117	0.1896
ALoRa	0.5265	0.5823	0.4902	0.3865	0.1394
TranAD	0.5660	0.5786	0.4919	0.3834	0.1284
TimesNet	0.3920	0.4233	0.4755	0.3598	0.1124
DLinear	0.4024	0.4795	0.4879	0.3695	0.1064
Autoformer	0.4687	0.5285	0.4985	0.3822	0.1107
Anomaly Transformer	0.5978	0.5069	0.2496	0.1571	0.0744
DeepSVDD	0.4782	0.4664	0.0291	0.0153	0.0274

VI. RESULTS

We organize the evaluation around six questions that map to operational intrusion detection requirements: low-noise detection, cross-domain behavior, parameter sensitivity, component contribution, real-time feasibility, and alarm auditability.

A. RQ1: Low-False-Positive Detection Performance

Attentive TCN-WGAN-GP achieves the best F1 score at the strict operating point while keeping the observed test benign FPR substantially lower than the strongest supervised baseline. The MLP baseline obtains similar recall at this point, but its test benign FPR is 0.1210, compared with 0.0036 for Attentive TCN-WGAN-GP. From an operator’s perspective, this difference is central: the proposed detector produces far fewer benign alarms while preserving a useful detection signal.

Table II shows that the advantage is preserved after adding modern deep anomaly-detection baselines. GANomaly reaches a competitive F1 score, but its observed benign test FPR is much larger than that of Attentive TCN-WGAN-GP. ALoRa, TimesNet, DLinear, Autoformer, Anomaly Transformer, TranAD, and DeepSVDD are less effective under this window-level CICIDS2017 protocol, suggesting that generic reconstruction, association, or representation objectives do not automatically solve low-FPR IDS calibration. This is the main security takeaway of RQ1: the ranking of detectors changes when benign-alarm cost is treated as a first-class metric.

[†] DLinear/Autoformer rows in Table III reuse the available fair-run external results from target FPR = 0.05 (same window-level protocol) as reference values.

TABLE III
WINDOW-LEVEL COMPARISON AT A HIGHER-RECALL OPERATING POINT (TARGET FPR 0.15).

Model	AUC	AP	F1	Recall	Precision	Test FPR
Attentive TCN-WGAN-GP	0.9464	0.9351	0.8783	0.9678	0.8040	0.1729
MLP	0.7537	0.7386	0.7140	0.7048	0.7234	0.1974
RF	0.6623	0.7126	0.6781	0.7117	0.6476	0.2837
OneClassSVM	0.8139	0.6506	0.6616	0.5810	0.7682	0.1284
GANomaly	0.7409	0.7706	0.6585	0.7237	0.6041	0.3474
LSTM-AE	0.4262	0.4201	0.5943	1.0000	0.4228	1.0000
LSTM-AD	0.5867	0.4705	0.5943	1.0000	0.4228	1.0000
TranAD	0.5660	0.5786	0.4849	0.4181	0.5771	0.2244
DLinear [†]	0.4024	0.4795	0.4879	0.3695	0.7102	0.1064
Autoformer [†]	0.4687	0.5285	0.4985	0.3822	0.7166	0.1107
IsolationForest	0.3961	0.4525	0.3603	0.2450	0.6810	0.0841
DeepSVDD	0.4782	0.4664	0.3462	0.2419	0.6091	0.1137

TABLE IV
SWAT FORMAL ANOMALY-DETECTION SPLIT AT TARGET FPR 0.05.

Model	AUC	AP	F1	Recall	Precision	Test FPR
Attentive TCN-WGAN-GP	0.9904	0.9977	0.9950	0.9900	1.0000	0.0000
OneClassSVM	0.9926	0.9982	0.9913	0.9827	1.0000	0.0000
TimesNet	0.9951	0.9973	0.9913	0.9837	0.9991	0.0030
DLinear	0.9945	0.9971	0.9909	0.9828	0.9993	0.0030
Autoformer	0.9907	0.9962	0.9909	0.9828	0.9993	0.0030
TranAD	0.9800	0.9953	0.9885	0.9800	0.9972	0.0090
LSTM-AE	0.9800	0.9953	0.9867	0.9800	0.9935	0.0209
GANomaly	0.9971	0.9993	0.9864	0.9945	0.9785	0.0716
IsolationForest	0.9861	0.9964	0.9768	0.9599	0.9943	0.0179
ALoRa	0.9794	0.9938	0.9822	0.9802	0.9842	0.0518
LSTM-AD	0.8680	0.9661	0.8988	0.8288	0.9817	0.0507
Anomaly Transformer	0.8661	0.9620	0.8869	0.8098	0.9803	0.0536
DeepSVDD	0.8485	0.9314	0.8636	0.9918	0.7647	1.0000

TABLE V
UNSW-NB15 OFFICIAL TRAIN/TEST SPLIT AT TARGET FPR 0.15.

Model	AUC	AP	F1	Recall	Precision	Test FPR
IsolationForest	0.8327	0.8430	0.8272	0.9856	0.7127	0.4926
TranAD	0.8467	0.8666	0.8178	0.9219	0.7349	0.4124
DeepSVDD	0.7334	0.7047	0.8177	0.9965	0.6932	0.5466
Attentive TCN-WGAN-GP	0.8017	0.7549	0.7928	0.9037	0.7062	0.4660
GANomaly	0.9703	0.9721	0.7613	0.9996	0.6147	0.7768
LSTM-AE	0.7964	0.8249	0.7205	0.7243	0.7168	0.3548
OneClassSVM	0.7576	0.8197	0.6178	0.5081	0.7879	0.1696
TimesNet [†]	0.7133	0.7806	0.2696	0.1579	0.9220	0.0166
DLinear [†]	0.7400	0.8003	0.2599	0.1512	0.8968	0.0153
Autoformer [†]	0.7963	0.8377	0.5353	0.3681	0.9855	0.0092
LSTM-AD	0.6723	0.6976	0.5546	0.4480	0.7280	0.2075
Anomaly Transformer	0.5845	0.6577	0.3684	0.2511	0.6919	0.1386
ALoRa	0.7243	0.7553	0.2180	0.1263	0.7979	0.0397

B. RQ2: Cross-Dataset Threshold Transfer

On SWaT, Attentive TCN-WGAN-GP remains in the top performance tier and obtains the strongest calibrated F1 under the formal anomaly-detection protocol. TimesNet and ALoRa also perform strongly, which supports the view that SWaT is a favorable setting for temporal reconstruction-style baselines. This result is important because it shows that the low-FPR behavior is not limited to CICIDS2017; the same benign temporal modeling idea also transfers to an industrial-control setting when the split matches one-class deployment assumptions.

[†] In Table V, TimesNet/DLinear/Autoformer rows are from available fair external runs at target FPR = 0.05 under the same preprocessing protocol.

TABLE VI
ABLATION AT TARGET FPR 0.15. ALL ROWS USE $W = 128$, STRIDE 16,
AND $\alpha = 0.24$.

Pooling	Loss	AUC	AP	F1	Recall
attn	WGAN-GP	0.9464	0.9351	0.8783	0.9678
mean	vanilla	0.9469	0.9289	0.8469	0.8940
mean	WGAN-GP	0.8253	0.8218	0.7543	0.7216
attn	vanilla	0.8367	0.6872	0.6891	0.7099

UNSW-NB15 is more challenging for threshold transfer. Attentive TCN-WGAN-GP remains competitive at the higher-recall operating point, reaching 0.7928 F1 and 0.9037 recall, but it does not dominate all baselines. Several methods obtain high recall by raising many benign alarms, as reflected by large observed test FPR values. TimesNet, Anomaly Transformer, and ALoRa achieve lower benign test FPRs, but at the cost of substantially lower recall. We therefore interpret UNSW-NB15 as a stress test showing that cross-dataset threshold calibration remains difficult under distribution shift. This result is useful for deployment: it warns that a threshold calibrated on one traffic regime should not be blindly reused on another without benign recalibration.

On TON_IoT, Attentive TCN-WGAN-GP achieves 0.5556 AUC and 0.1794 F1 at target FPR 0.05. This is not a strong result, and we report it as a limitation rather than as evidence of universal superiority. The dataset uses telemetry-style features whose temporal dynamics differ from flow and process-control traffic, suggesting that domain adaptation and more robust threshold calibration are promising directions for future work.

C. RQ3: Hyperparameter Sensitivity

The final configuration uses $W = 128$, stride $r = 16$, fused scoring, and $\alpha = 0.24$. The staged selection protocol prevents the ablation from changing multiple experimental factors simultaneously.

The selected α indicates that the final score relies more on feature-deviation evidence than on the critic score alone. This is consistent with the intuition that adversarial training provides useful representations, but deployment-time anomaly scoring can benefit from comparing these representations with benign reference statistics.

D. RQ4: Ablation Study

The ablation shows that attention alone is not sufficient: the attention-vanilla variant is weaker than the mean-vanilla baseline. However, attention combined with WGAN-GP achieves the best F1 and recall at the higher-recall operating point. A five-seed check further shows that strict low-FPR calibration is more sensitive to random initialization.

E. RQ5: Real-Time Monitoring Feasibility

Table VII reports eval-only throughput for saved checkpoints using fused scoring. The final CICIDS2017 checkpoint

TABLE VII
INFERENCE EFFICIENCY OF ATTENTIVE TCN-WGAN-GP CHECKPOINTS.

Dataset	Windows	Seconds	Windows/s	MB
CICIDS2017	43897	10.51	4176.3	1.56
SWaT	1433	8.09	177.1	1.53
UNSW-NB15	5138	31.48	163.2	1.46
TON_IoT	18743	70.67	265.2	1.39

TABLE VIII
LIGHTWEIGHT GAUSSIAN-NOISE ROBUSTNESS ON 12,000 SAMPLED
CICIDS2017 TEST WINDOWS. THE THRESHOLD IS CALIBRATED ON
CLEAN BENIGN WINDOWS AT TARGET FPR 0.05.

Scenario	F1	Recall	Precision	Test FPR
Clean	0.7552	0.6945	0.8276	0.1447
Noise 1%	0.7596	0.6967	0.8350	0.1377
Noise 3%	0.7710	0.7047	0.8512	0.1232
Noise 5%	0.7866	0.7248	0.8598	0.1182

evaluates tens of thousands of windows in seconds, and all checkpoints remain compact at about 1.4–1.6 MB. This supports offline real-time feasibility at the window level.

Table VIII gives a lightweight robustness check under small Gaussian perturbations after min-max scaling. The sampled-window results do not show degradation under 1–5% noise, but they should be interpreted as a sanity check rather than a full adversarial robustness guarantee.

Because the detector operates on overlapping windows, the practical throughput depends on both model inference speed and stride. With stride $r = 16$, the detector does not rescore every individual flow as a completely independent decision; instead, it emits one score per window step. This design reduces alert jitter and creates a controllable latency-throughput trade-off. In a streaming deployment, the model can maintain a rolling buffer of the latest W flow vectors and update the score whenever r new flows arrive. The resulting latency is measured in newly observed flows rather than in full retraining cycles, which is appropriate for a monitoring component placed behind a flow collector.

F. RQ6: Alarm Auditability

The XAI analysis indicates that the final model bases anomaly scores on interpretable traffic features such as TCP control flags, packet rates, header-length statistics, destination ports, and packet-length variation. The complete artifact package includes feature attribution, temporal attribution, discriminator attention-weight plots, and a single-window case study. These explanations help analysts inspect not only whether a window is anomalous, but also when and through which traffic characteristics the model raises an alarm. We treat these outputs as alarm-audit artifacts: they support triage and model debugging, but they do not replace packet-level forensic investigation.




Fig. 3. Single-window multi-view XAI case study. The anomalous case has score 1.0000, while a benign reference case has score 0.0506. The figure compares attention, Integrated Gradients, SHAP, and top feature evidence.

TABLE IX

TOP-FEATURE STABILITY OVER 64 ANOMALOUS CICIDS2017 WINDOWS USING GRADIENT-TIMES-INPUT ATTRIBUTION. COVERAGE IS THE FRACTION OF WINDOWS IN WHICH A FEATURE APPEARS IN THE WINDOW’S TOP-5 ATTRIBUTED FEATURES.

Feature	Top-5 Count	Coverage	Mean Rank
PSH Flag Count	53	0.828	1.23
Flow Packets/s	42	0.656	2.74
Fwd Header Length.1	40	0.625	3.90
min_seg_size_forward	39	0.609	4.21
Destination Port	33	0.516	3.45

The distinction between attention and attribution is important for interpreting the model. Attention weights show how the discriminator pools temporal evidence before producing a critic score. Gradient-times-input attributions, in contrast, estimate how sensitive the final fused anomaly score is to each time-feature entry.

Fig. 3 shows a concrete alarm case with attention, Integrated Gradients, SHAP, and feature evidence. The most important features include TCP control flags, packet-rate statistics, header-length statistics, destination ports, and backward packet-length statistics. Table IX further shows that several feature groups recur across many anomalous windows rather than appearing only in one hand-picked example. The case study also illustrates why attention is not used alone: attention highlights temporal aggregation behavior, while Integrated Gradients and SHAP provide post-hoc evidence for the final anomaly score.

VII. DISCUSSION

The staged selection protocol is a pragmatic design choice. It makes the final comparison easier to interpret because the ablation changes only the model components under study, rather than changing window size, stride, and score fusion simultaneously.

The cross-dataset results refine the scope of the claim. CICIDS2017 is the primary benchmark on which the proposed detector is strongest, and SWaT provides a meaningful transfer check in an industrial-control setting. UNSW-NB15 should be interpreted as a stress test rather than as a second main win condition: the detector remains competitive, but the ranking among anomaly detectors changes under distribution shift and several methods trade substantially higher benign alarms for recall. TON_IoT is more difficult still and is included as an explicit out-of-domain limitation. We therefore do not claim universal dominance across all network and telemetry domains. Instead, we claim that the proposed detector is particularly effective in the target low-FPR window-level IDS setting and that threshold transfer remains an important open problem.

For a deployable network security system, the most important lesson is that model accuracy and alarm policy should be evaluated together. Several baselines can obtain high recall only by increasing benign alarms, while strict calibration can reduce analyst burden but may miss attacks. The proposed framework makes this trade-off explicit by reporting both requested calibration budgets and observed test FPR. This also suggests a practical deployment path: calibrate thresholds on recent benign traffic from the target environment, monitor benign-alarm drift, and periodically refresh calibration rather than assuming that a fixed benchmark threshold will transfer unchanged.

The current XAI module should be interpreted as post-hoc analysis. It helps inspect whether model alarms are driven by stable temporal and feature-level evidence, but the current paper does not claim that the XAI signal improves training or provides a formal faithfulness guarantee. A future extension could use attribution-guided regularization, feature selection, or domain-adaptive calibration, but those should be evaluated as separate model contributions.

VIII. ETHICAL CONSIDERATIONS

This work uses public benchmark network traffic data and does not run attacks against real systems. The primary risk is misuse or overclaiming: an intrusion detector trained on a benchmark may not generalize to all operational networks. We therefore report false positive rates, throughput, and limitations, and avoid presenting the model as a standalone replacement for human security operations. In deployment, the detector should be used as a triage aid whose alarms are reviewed together with logs, packet evidence, and existing security controls.

IX. OPEN SCIENCE AND ARTIFACT AVAILABILITY

The codebase contains scripts for staged tuning, baseline comparison, ablation, FPR-sweep analysis, seed runs, and XAI generation. For submission, the artifact should include anonymized source code, configuration files, plotting scripts, and instructions to reproduce the reported tables and figures.

X. CONCLUSION

This paper presents Attentive TCN-WGAN-GP for low-false-positive, window-level network intrusion detection. The model combines temporal convolutional sequence modeling, adversarial benign-window learning, attention pooling, WGAN-GP training, fused anomaly scoring, and post-hoc XAI analysis. On CICIDS2017, the final model outperforms same-granularity and modern anomaly-detection baselines at the strict low-FPR operating point and achieves a stronger recall/F1 trade-off when the operating point is relaxed. On SWaT, it preserves a strong low-FPR operating point in a distinct industrial-control setting. UNSW-NB15 shows that the method remains competitive under distribution shift, but also clarifies that threshold transfer is difficult and that performance rankings change when benign-alarm cost is made explicit. TON_IoT further marks the boundary of the method's current cross-domain robustness. The ablation and multi-seed results show that attention should be paired with WGAN-GP and that strict tail-threshold calibration introduces variance. The XAI case study and cross-window feature-stability analysis further show that alarms can be audited through meaningful temporal and feature-level evidence. Overall, the results support a cautious but operationally relevant claim: temporal benign-window modeling is a strong basis for low-noise NIDS alarms in the target setting, but deployment still requires environment-specific calibration and analyst-facing evidence.

REFERENCES

- [1] S. Bai, J. Z. Kolter, and V. Koltun, "An empirical evaluation of generic convolutional and recurrent networks for sequence modeling," *arXiv preprint arXiv:1803.01271*, 2018.
- [2] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial nets," in *Advances in Neural Information Processing Systems*, 2014.
- [3] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein generative adversarial networks," in *International Conference on Machine Learning*, 2017.
- [4] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville, "Improved training of wasserstein gans," in *Advances in Neural Information Processing Systems*, 2017.
- [5] M. Tavallaei, E. Bagheri, W. Lu, and A. A. Ghorbani, "A detailed analysis of the kdd cup 99 data set," in *IEEE Symposium on Computational Intelligence for Security and Defense Applications*, 2009.
- [6] N. Moustafa and J. Slay, "Unsw-nb15: A comprehensive data set for network intrusion detection systems," in *Military Communications and Information Systems Conference*, 2015.
- [7] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *International Conference on Information Systems Security and Privacy*, 2018.
- [8] F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation forest," in *IEEE International Conference on Data Mining*, 2008.
- [9] B. Scholkopf, J. C. Platt, J. Shawe-Taylor, A. J. Smola, and R. C. Williamson, "Estimating the support of a high-dimensional distribution," *Neural Computation*, vol. 13, no. 7, pp. 1443–1471, 2001.
- [10] Y. Mirsky, T. Doitshman, Y. Elovici, and A. Shabtai, "Kitsune: An ensemble of autoencoders for online network intrusion detection," in *Network and Distributed System Security Symposium*, 2018.
- [11] L. D. Manocchio, S. Layeghy, W. W. Lo, G. K. Kulatilleke, M. Sarhan, and M. Portmann, "Flowtransformer: A transformer framework for flow-based network intrusion detection systems," *Expert Systems with Applications*, vol. 241, p. 122564, 2024.
- [12] Z. Long, H. Yan, G. Shen, X. Zhang, H. He, and L. Cheng, "A transformer-based network intrusion detection approach for cloud security," *Journal of Cloud Computing*, vol. 13, no. 5, 2024.
- [13] C. Xi, H. Wang, and X. Wang, "A novel multi-scale network intrusion detection model with transformer," *Scientific Reports*, vol. 14, p. 23239, 2024.
- [14] P. F. d. Araujo-Filho, M. Naili, G. Kaddoum, E. T. Fapi, and Z. Zhu, "Unsupervised gan-based intrusion detection system using temporal convolutional networks and self-attention," *IEEE Transactions on Network and Service Management*, vol. 20, no. 4, pp. 4951–4963, 2023.
- [15] V. G. d. S. Ruffo, D. M. B. Lent, L. F. Carvalho, J. Lloret, and M. L. Proenca Jr., "Generative adversarial networks to detect intrusion and anomaly in ip flow-based networks," *Future Generation Computer Systems*, vol. 163, p. 107531, 2025.
- [16] P. Malhotra, A. Ramakrishnan, G. Anand, L. Vig, P. Agarwal, and G. Shroff, "Lstm-based encoder-decoder for multi-sensor anomaly detection," in *ICML Workshop on Anomaly Detection*, 2016.
- [17] M. Du, F. Li, G. Zheng, and V. Srikumar, "Deeplog: Anomaly detection and diagnosis from system logs through deep learning," in *ACM Conference on Computer and Communications Security*, 2017.
- [18] M. Sakurada and T. Yairi, "Anomaly detection using autoencoders with nonlinear dimensionality reduction," in *Workshop on Machine Learning for Sensory Data Analysis*, 2014.
- [19] J. An and S. Cho, "Variational autoencoder based anomaly detection using reconstruction probability," *Special Lecture on IE*, vol. 2, no. 1, pp. 1–18, 2015.
- [20] L. Ruff, R. A. Vandermeulen, N. Goernitz, L. Deecke, S. A. Siddiqui, A. Binder, E. Mueller, and M. Kloft, "Deep one-class classification," in *International Conference on Machine Learning*, 2018, pp. 4393–4402.
- [21] S. Tuli, G. Casale, and N. R. Jennings, "Tranad: Deep transformer networks for anomaly detection in multivariate time series data," *Proceedings of the VLDB Endowment*, vol. 15, no. 6, pp. 1201–1214, 2022.
- [22] J. Xu, H. Wu, J. Wang, and M. Long, "Anomaly transformer: Time series anomaly detection with association discrepancy," in *International Conference on Learning Representations*, 2022.
- [23] H. Wu, T. Hu, Y. Liu, H. Zhou, J. Wang, and M. Long, "Timesnet: Temporal 2d-variation modeling for general time series analysis," in *International Conference on Learning Representations*, 2023.
- [24] M. Landauer, F. Skopik, B. Stojanovic, A. Flatscher, and T. Ullrich, "A review of time-series analysis for cyber security analytics: From intrusion detection to attack prediction," *International Journal of Information Security*, vol. 24, no. 3, 2025.
- [25] Y. Kong, J. Guo, L. Ren, G. Kang, Y. Wang, and J. Lu, "Dadn: A dynamic anomaly detection network for multivariate time series data of the industrial internet of things," *IEEE Transactions on Industrial Informatics*, vol. 21, no. 8, pp. 6294–6304, 2025.
- [26] T. Schlegl, P. Seebock, S. M. Waldstein, U. Schmidt-Erfurth, and G. Langs, "Unsupervised anomaly detection with generative adversarial networks to guide marker discovery," in *International Conference on Information Processing in Medical Imaging*, 2017.
- [27] T. Schlegl, P. Seebock, S. M. Waldstein, G. Langs, and U. Schmidt-Erfurth, "f-anogan: Fast unsupervised anomaly detection with generative adversarial networks," *Medical Image Analysis*, vol. 54, pp. 30–44, 2019.
- [28] S. Akcay, A. Atapour-Abarghouei, and T. P. Breckon, "Ganomaly: Semi-supervised anomaly detection via adversarial training," in *Asian Conference on Computer Vision*, 2018, pp. 622–637.
- [29] K. Simonyan, A. Vedaldi, and A. Zisserman, "Deep inside convolutional networks: Visualising image classification models and saliency maps," in *International Conference on Learning Representations Workshop*, 2014.
- [30] M. Sundararajan, A. Taly, and Q. Yan, "Axiomatic attribution for deep networks," in *International Conference on Machine Learning*, 2017.
- [31] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should i trust you? explaining the predictions of any classifier," in *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016.
- [32] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems*, 2017.
- [33] M. A. Rahman, G. Francia, H. Shahriar, E. El-Sheikh, A. Cuzzocrea, and S. I. Ahamed, "Explaining network intrusion detection system with shap and lime," in *IEEE International Conference on Big Data*, 2025, pp. 4325–4332.
- [34] J. Sehgal, R. Kumar, and R. K. Bedi, "Adaptive and explainable slwambilstm intrusion detection system with drift awareness and cross-dataset evaluation," in *International Conference on Mobile Computing and Sustainable Informatics*, 2026, pp. 956–962.

- [35] S. Jain and B. C. Wallace, "Attention is not explanation," in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, 2019.
- [36] J. Goh, S. Adepu, K. N. Junejo, and A. Mathur, "A dataset to support research in the design of secure water treatment systems," in *International Conference on Critical Information Infrastructures Security*, 2016, pp. 88–99.
- [37] M. M. Aslam, A. Tufail, L. C. De Silva, R. A. A. H. M. Apong, and A. Namoun, "An improved autoencoder-based approach for anomaly detection in industrial control systems," *Systems Science & Control Engineering*, 2024.
- [38] A. Alsaedi, N. Moustafa, Z. Tari, A. Mahmood, and A. Anwar, "Ton_iot telemetry dataset: A new generation dataset of iot and iiot for data-driven intrusion detection systems," *IEEE Access*, vol. 8, pp. 165 130–165 150, 2020.
- [39] A. A. Chavan, J. Ingale, S. Pimpalkar, and S. Pawar, "A hybrid autoencoder-xgboost model for network intrusion detection: Performance analysis on unsw-nb15 and ton_iot network datasets," in *Control Instrumentation System Conference*, 2025, pp. 1–6.
- [40] ALoRa Authors, "Alora: Adaptive localization and representation for time-series anomaly detection," <https://github.com/CharisShimillas/ALoRa>, 2024.
- [41] A. Zeng, M. Chen, L. Zhang, and Q. Xu, "Are transformers effective for time series forecasting?" in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2023.
- [42] H. Wu, J. Xu, J. Wang, and M. Long, "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting," in *Advances in Neural Information Processing Systems*, 2021.
- [43] Y. Wang, H. Wu, J. Dong, Y. Liu, M. Long, and J. Wang, "Deep time series models: A comprehensive survey and benchmark," *arXiv preprint arXiv:2407.13278*, 2024.